

# incomeprediction

March 24, 2023

## 1 Importing all necessary libraries

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

## 2 Data Ingestion

```
[2]: df=pd.read_csv(r'C:\Users\PS4Z\Downloads\archive\income_evaluation.csv')
```

```
[3]: #seeing how the looks like
df.head()
```

```
[3]:
```

	age	workclass	fnlwgt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	marital-status	occupation	relationship	race	sex	\
0	Never-married	Adm-clerical	Not-in-family	White	Male	
1	Married-civ-spouse	Exec-managerial	Husband	White	Male	
2	Divorced	Handlers-cleaners	Not-in-family	White	Male	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female	

	capital-gain	capital-loss	hours-per-week	native-country	income
0	2174	0	40	United-States	<=50K
1	0	0	13	United-States	<=50K
2	0	0	40	United-States	<=50K
3	0	0	40	United-States	<=50K
4	0	0	40	Cuba	<=50K

### 3 Understanding data

```
[4]: #seeing the shape of data
print('Data Shape: ',df.shape)
```

Data Shape: (32561, 15)

```
[5]: #understanding about null values in data
df.isnull().sum()
```

```
[5]: age                0
     workclass          0
     fnlwgt             0
     education          0
     education-num      0
     marital-status     0
     occupation         0
     relationship       0
     race               0
     sex                0
     capital-gain       0
     capital-loss       0
     hours-per-week     0
     native-country     0
     income             0
dtype: int64
```

```
[6]: #Getting information about data; null counts and data types of data columns
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   age                   32561 non-null  int64
1   workclass             32561 non-null  object
2   fnlwgt                32561 non-null  int64
3   education             32561 non-null  object
4   education-num         32561 non-null  int64
5   marital-status        32561 non-null  object
6   occupation            32561 non-null  object
7   relationship          32561 non-null  object
8   race                  32561 non-null  object
9   sex                   32561 non-null  object
10  capital-gain           32561 non-null  int64
11  capital-loss           32561 non-null  int64
12  hours-per-week         32561 non-null  int64
```

```
13  native-country  32561 non-null  object
14  income          32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

Observations: There are total 32537 rows and 15 columns in the dataset Categorical features = 9 and Numerical features = 6

```
[7]: #list of column names
df.columns
```

```
[7]: Index(['age', 'workclass', 'fnlwgt', 'education', 'education-num',
         'marital-status', 'occupation', 'relationship', 'race', 'sex',
         'capital-gain', 'capital-loss', 'hours-per-week', 'native-country',
         'income'],
        dtype='object')
```

```
[8]: #getting rid of all spaces in the column names
df.columns= df.columns.str.strip()
```

```
[9]: df.columns
```

```
[9]: Index(['age', 'workclass', 'fnlwgt', 'education', 'education-num',
         'marital-status', 'occupation', 'relationship', 'race', 'sex',
         'capital-gain', 'capital-loss', 'hours-per-week', 'native-country',
         'income'],
        dtype='object')
```

```
[10]: #getting data types of each column header
df.dtypes
```

```
[10]: age                int64
workclass             object
fnlwgt                int64
education             object
education-num         int64
marital-status        object
occupation            object
relationship          object
race                 object
sex                  object
capital-gain          int64
capital-loss          int64
hours-per-week        int64
native-country        object
income               object
dtype: object
```

```
[11]: #getting a closer look on our target variable
df['income'].value_counts()
```

```
[11]: <=50K    24720
      >50K     7841
      Name: income, dtype: int64
```

Observation: Most employees fall in <=50K income category

```
[12]: #classifying our target variable in binary notations for ease of understanding
df['income']=df['income'].map({' <=50K':0, ' >50K':1})
```

```
[13]: #seeing the data type of target variable change from object to int
df.dtypes
```

```
[13]: age                int64
      workclass         object
      fnlwgt           int64
      education         object
      education-num     int64
      marital-status    object
      occupation        object
      relationship      object
      race              object
      sex               object
      capital-gain      int64
      capital-loss      int64
      hours-per-week    int64
      native-country    object
      income            int64
      dtype: object
```

```
[14]: #duplicate entries in data
df.duplicated().sum()
```

```
[14]: 24
```

```
[15]: #dropping all duplicate entries
df.drop_duplicates(inplace=True)
```

```
[16]: df.head()
```

```
[16]:   age  workclass  fnlwgt  education  education-num  \
0   39   State-gov   77516   Bachelors             13
1   50  Self-emp-not-inc   83311   Bachelors             13
2   38    Private  215646   HS-grad              9
3   53    Private  234721    11th              7
```

4    28                    Private   338409   Bachelors                    13

	marital-status	occupation	relationship	race	sex \
0	Never-married	Adm-clerical	Not-in-family	White	Male
1	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female

	capital-gain	capital-loss	hours-per-week	native-country	income
0	2174	0	40	United-States	0
1	0	0	13	United-States	0
2	0	0	40	United-States	0
3	0	0	40	United-States	0
4	0	0	40	Cuba	0

```
[17]: #getting value counts for workclass column
df['workclass'].value_counts()
```

```
[17]: Private                22673
Self-emp-not-inc          2540
Local-gov                 2093
?                          1836
State-gov                 1298
Self-emp-inc              1116
Federal-gov                960
Without-pay               14
Never-worked               7
Name: workclass, dtype: int64
```

```
[18]: #getting value counts for native-country column
df['native-country'].value_counts()
```

```
[18]: United-States          29153
Mexico                     639
?                           582
Philippines                198
Germany                   137
Canada                    121
Puerto-Rico               114
El-Salvador               106
India                     100
Cuba                       95
England                   90
Jamaica                   81
South                     80
China                     75
```

Italy	73
Dominican-Republic	70
Vietnam	67
Japan	62
Guatemala	62
Poland	60
Columbia	59
Taiwan	51
Haiti	44
Iran	43
Portugal	37
Nicaragua	34
Peru	31
France	29
Greece	29
Ecuador	28
Ireland	24
Hong	20
Cambodia	19
Trinidad&Tobago	19
Laos	18
Thailand	18
Yugoslavia	16
Outlying-US(Guam-USVI-etc)	14
Honduras	13
Hungary	13
Scotland	12
Holand-Netherlands	1

Name: native-country, dtype: int64

```
[19]: #getting value counts for occupation column
df['occupation'].value_counts()
```


```
[19]: Prof-specialty      4136
      Craft-repair      4094
      Exec-managerial    4065
      Adm-clerical       3768
      Sales              3650
      Other-service      3291
      Machine-op-inspct  2000
      ?                  1843
      Transport-moving   1597
      Handlers-cleaners  1369
      Farming-fishing     992
      Tech-support        927
      Protective-serv     649
      Priv-house-serv     147
```

```
Armed-Forces          9
Name: occupation, dtype: int64
```

Observation:workclass, occupation and native\_country has missing values

## 4 Understanding Imbalanced Data

```
[20]: #Plotting Bar Plot for target variable
df['income'].value_counts(normalize=True).plot(kind='bar');
```



Observation:we can see above data is imbalanced, more class belongs to less than or equal to 50k

```
[21]: #getting a sample from data
df.sample(7)
```

```
[21]:
```

	age	workclass	fnlwgt	education	education-num	\
15765	29	Private	94892	Some-college	10	
29330	37	Private	212437	Some-college	10	
19533	26	Private	86483	Some-college	10	
20227	40	?	65545	Masters	14	
14230	62	Local-gov	33365	HS-grad	9	
13385	54	Self-emp-inc	383365	Bachelors	13	

20389	50	Self-emp-not-inc	30731	Assoc-voc	11
-------	----	------------------	-------	-----------	----

	marital-status	occupation	relationship	race	\
15765	Married-civ-spouse	Prof-specialty	Husband	White	
29330	Widowed	Machine-op-inspct	Unmarried	Black	
19533	Never-married	Exec-managerial	Not-in-family	White	
20227	Divorced	?	Own-child	White	
14230	Widowed	Other-service	Not-in-family	White	
13385	Married-civ-spouse	Sales	Husband	White	
20389	Never-married	Other-service	Not-in-family	White	

	sex	capital-gain	capital-loss	hours-per-week	native-country	\
15765	Male	0	0	40	United-States	
29330	Female	0	0	48	United-States	
19533	Female	0	0	40	United-States	
20227	Female	0	0	55	United-States	
14230	Female	0	0	40	Canada	
13385	Male	0	0	70	United-States	
20389	Male	0	0	50	United-States	

	income
15765	0
29330	0
19533	0
20227	0
14230	0
13385	1
20389	0

```
[22]: #Getting 5 point summary for all numercial features
df.describe()
```

```
[22]:
```

	age	fnlwgt	education-num	capital-gain	capital-loss	\
count	32537.000000	3.253700e+04	32537.000000	32537.000000	32537.000000	
mean	38.585549	1.897808e+05	10.081815	1078.443741	87.368227	
std	13.637984	1.055565e+05	2.571633	7387.957424	403.101833	
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	
75%	48.000000	2.369930e+05	12.000000	0.000000	0.000000	
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	

	hours-per-week	income
count	32537.000000	32537.000000
mean	40.440329	0.240926
std	12.346889	0.427652
min	1.000000	0.000000



25%	40.000000	0.000000
50%	40.000000	0.000000
75%	45.000000	0.000000
max	99.000000	1.000000

## 5 Separating numerical features from categorical features

```
[23]: #categorizing numericl and categorical features based on the data type
num_feat=[i for i in df.columns if df[i].dtypes!='O']
cat_feat=[j for j in df.columns if df[j].dtypes=='O']
```

```
[24]: num_feat, cat_feat
```


```
[24]: ([ 'age',
        'fnlwgt',
        'education-num',
        'capital-gain',
        'capital-loss',
        'hours-per-week',
        'income'],
       [ 'workclass',
        'education',
        'marital-status',
        'occupation',
        'relationship',
        'race',
        'sex',
        'native-country'])
```

## 6 Visualizing the data

```
[25]: #pairplot shows graphical representation of all numerical features with one_
      ↪ another with target variable in legend
sns.pairplot(data=df,hue='income');
```

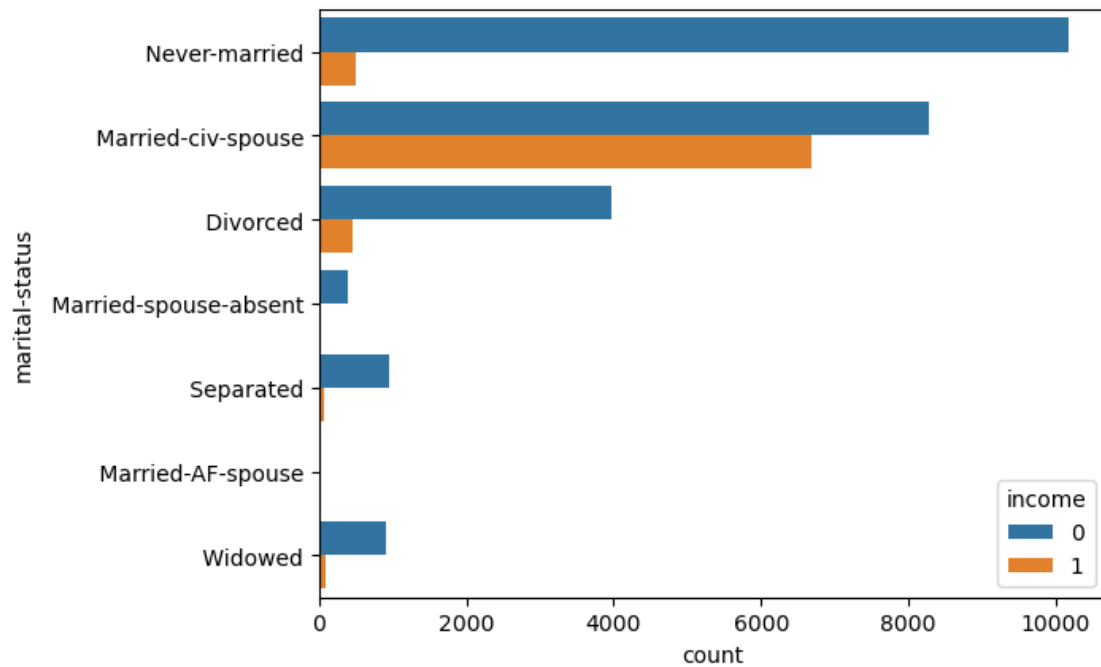


```
[26]: #countplot of sex of employees with legend
sns.countplot(y=df['sex'],hue=df['income']);
```




Observation:from above plot, we can say that male are earning more than female

```
[27]: #countplot of marital Status of employees with legend  
sns.countplot(y=df['marital-status'],hue=df['income']);
```



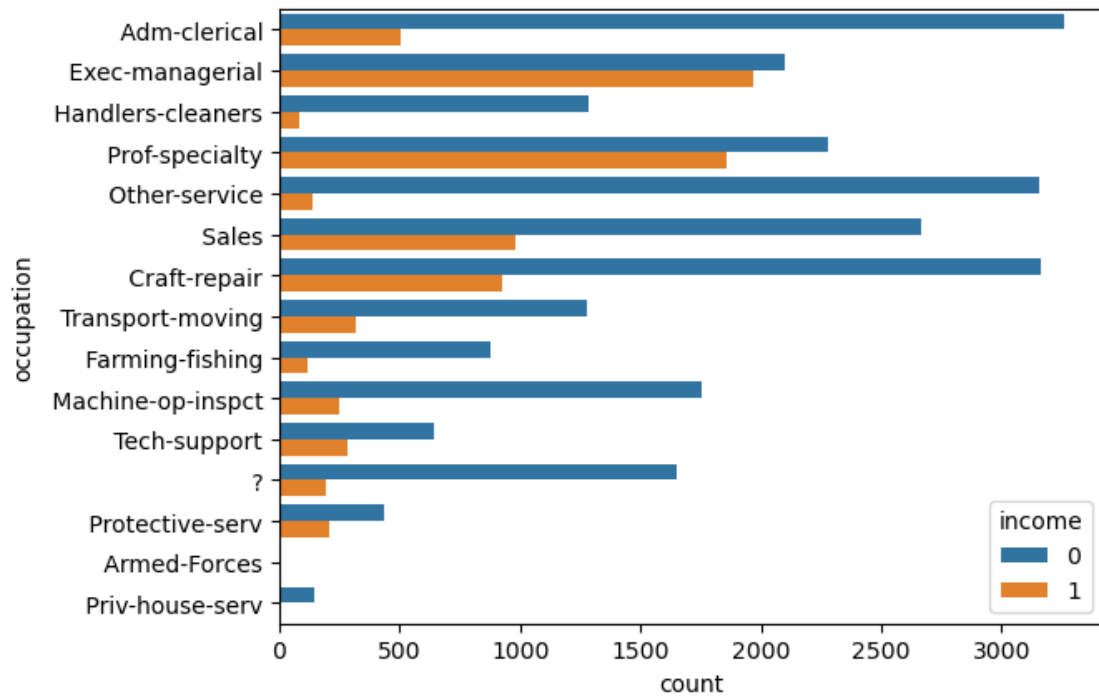
Observation: people belonging to 'never married' status, are earning more than that of people belonging to other marital status

```
[28]: #countplot of relationship of employees with legend
sns.countplot(y=df['relationship'], hue=df['income']);
```




Observation: people belonging to 'Not-in-family' and 'Husband' status, are earning more than that of people belonging to other relationship status.

```
[29]: #countplot of occupation of employees with legend  
sns.countplot(y=df['occupation'], hue=df['income']);
```




Observation: people belonging to 'Adm-clerical' occupation, are earning more than that of people belonging to other occupations.

```
[30]: #countplot of education of employees with legend
sns.countplot(y=df['education'], hue=df['income']);
```



Observation: people belonging to 'HS-grad' status, are earning more than that of people belonging to other education qualification


```
[31]: #countplot of race of employees with legend
sns.countplot(y=df['race'], hue=df['income']);
```



Observation: people belonging to 'white' race, are earning more than that of people belonging to other races


```
[32]: #Bivariate analysis of numerical features w.r.t Target Variable
plt.figure(figsize=(10,10))
plt.suptitle('Analysis w.r.t Target Variable')
for a in range(0,len(num_feat)):
    plt.subplot(3,3,a+1)
    sns.barplot(x='income',y=num_feat[a],data=df)
plt.tight_layout()
```





Observations: The bar plot above are bivariate plots. In terms of Age, older the person, more is the probability of income getting higher. Capital gain, and capital loss are more experienced by people having higher income.


```
[33]: #barplot of sex with legend of Target Variable
sns.barplot(x=df['sex'],y=df['hours-per-week'],hue=df['income'],data=df);
```




Observations:male works more hours than female

## 7 Checking for Special Symbol


```
[34]: df['workclass'].value_counts().plot(kind='barh');
```



```
[35]: df['occupation'].value_counts().plot(kind='barh');
```



```
[36]: df['native-country'].value_counts().plot(kind='bar')
plt.tight_layout();
```



Observations: ' ?' symbol in workclass, occupation and native-country are most possibly NaN values

## 8 Replacing the Special Symbol

```
[37]: df.replace(' ?',np.NaN,inplace=True)
```


```
[38]: df.isna().sum()
```

```
[38]: age          0
workclass      1836
fnlwgt         0
education       0
education-num   0
marital-status  0
occupation     1843
relationship    0
```


```
race          0
sex           0
capital-gain  0
capital-loss  0
hours-per-week 0
native-country 582
income        0
dtype: int64
```

```
[39]: for col in ['workclass', 'occupation', 'native-country']:
      df[col].fillna(df[col].mode()[0], inplace=True)
```


```
[40]: df['workclass'].value_counts().plot(kind='barh');
```



```
[41]: df['occupation'].value_counts().plot(kind='barh');
```



```
[42]: df['native-country'].value_counts().plot(kind='bar');
```



Observations: NaN values are replaced and we can see in above three plot the value is increased as compared to the previous plots

```
[43]: df.corr()
```

```
[43]:
```


	age	fnlwgt	education-num	capital-gain	capital-loss	\
age	1.000000	-0.076447	0.036224	0.077676	0.057745	
fnlwgt	-0.076447	1.000000	-0.043388	0.000429	-0.010260	
education-num	0.036224	-0.043388	1.000000	0.122664	0.079892	
capital-gain	0.077676	0.000429	0.122664	1.000000	-0.031639	
capital-loss	0.057745	-0.010260	0.079892	-0.031639	1.000000	
hours-per-week	0.068515	-0.018898	0.148422	0.078408	0.054229	

income	0.234037	-0.009502	0.335272	0.223336	0.150501
--------	----------	-----------	----------	----------	----------

	hours-per-week	income
age	0.068515	0.234037
fnlwgt	-0.018898	-0.009502
education-num	0.148422	0.335272
capital-gain	0.078408	0.223336
capital-loss	0.054229	0.150501
hours-per-week	1.000000	0.229658
income	0.229658	1.000000


```
[44]: sns.heatmap(data=df.corr(),annot=True,cmap='viridis');
```



Observations: there is very less correlation between features, which shows there is no multicollinearity and fnlwgt has the lowest correlation with the rest of the features; the highest correlation is between education and income.

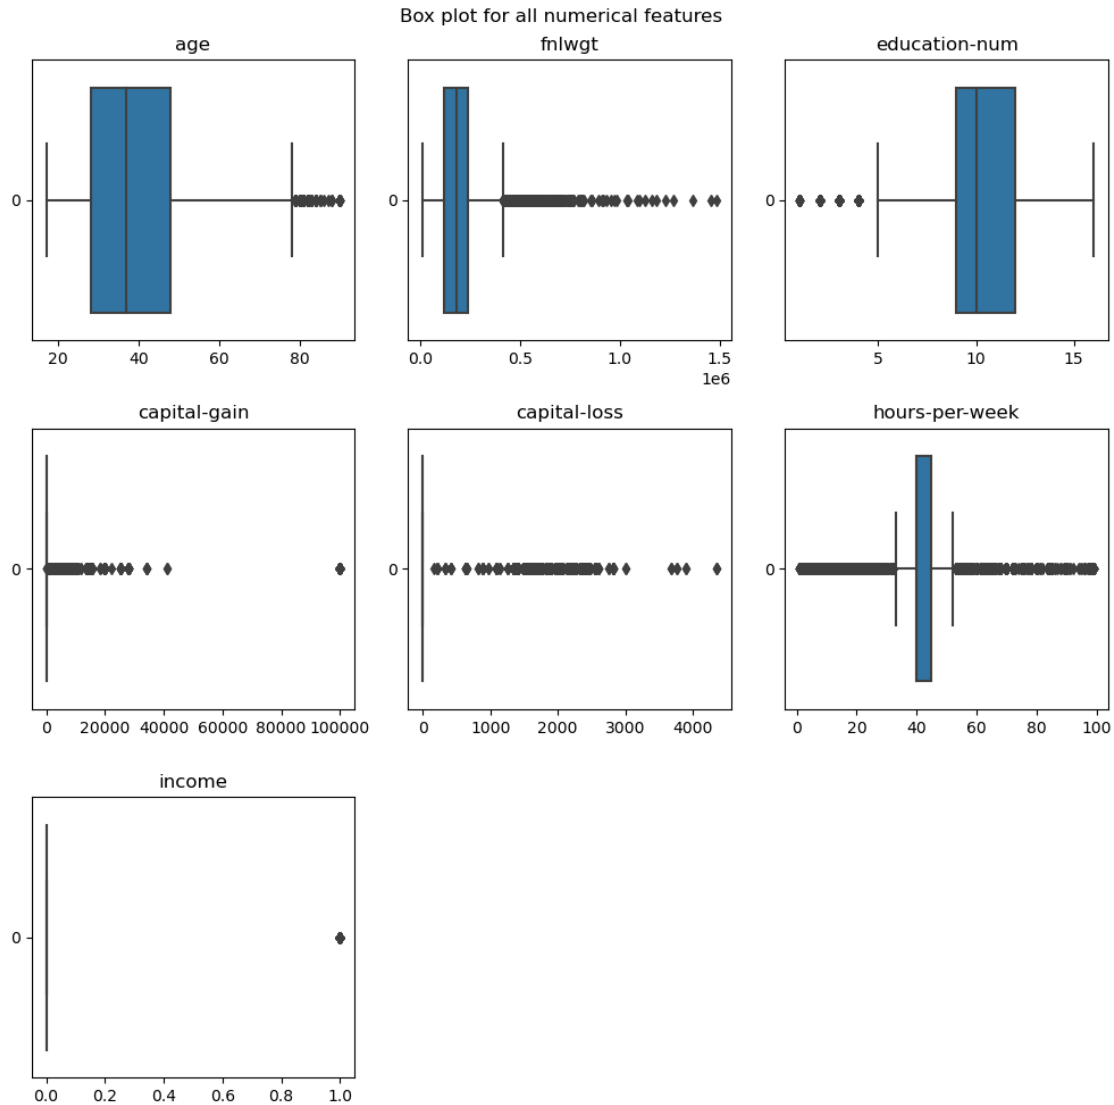


```
[45]: plt.figure(figsize=(10,5))
sns.countplot(x=df['workclass'],hue=df['income'])
plt.tight_layout();
```



Observations:employees in ‘Private’ workclass get paid the most and within those, most of them falls in <=50K pay scale

```
[46]: #Bivariate analysis of numerical features w.r.t Target Variable
plt.figure(figsize=(10,10))
plt.suptitle('Box plot for all numerical features')
for a in range(0,len(num_feat)):
    plt.subplot(3,3,a+1)
    sns.boxplot(data=df[num_feat[a]],orient='h')
    plt.title(label=num_feat[a])
plt.tight_layout()
```




## 9 Handling Outliers

```
[47]: #deriving upperlimit and lowerlimit for 'fnlwgt'
IQR= df['fnlwgt'].quantile(0.75)-df['fnlwgt'].quantile(0.25)
Lower=df['fnlwgt'].quantile(0.25)-1.5*IQR
Upper=df['fnlwgt'].quantile(0.75)+1.5*IQR
print('Lower_limit:',Lower)
print('Upper_limit:',Upper)
```

```
Lower_limit: -60922.0
Upper_limit: 415742.0
```

```
[48]: #getting rid of outliers
df['fnlwgt']=np.where(df['fnlwgt']>Upper,Upper,np.
    ↳where(df['fnlwgt']<Lower,Lower,df['fnlwgt']))
```

```
[49]: #boxplot must not show outliers now
sns.boxplot(x='fnlwgt', data=df);
```




```
[50]: #deriving upperlimit and lowerlimit for 'age'
IQR= df['age'].quantile(0.75)-df['age'].quantile(0.25)
Lower=df['age'].quantile(0.25)-1.5*IQR
Upper=df['age'].quantile(0.75)+1.5*IQR
print('Lower_limit:',Lower)
print('Upper_limit:',Upper)
```

```
Lower_limit: -2.0
Upper_limit: 78.0
```

```
[51]: #getting rid of outliers
df['age']=np.where(df['age']>Upper,Upper,np.
    ↳where(df['age']<Lower,Lower,df['age']))
```

```
[52]: #boxplot must not show outliers now
sns.boxplot(x='age', data=df);
```




```
[53]: #deriving upperlimit and lowerlimit for 'hours-per-week'
IQR= df['hours-per-week'].quantile(0.75)-df['hours-per-week'].quantile(0.25)
Lower=df['hours-per-week'].quantile(0.25)-1.5*IQR
Upper=df['hours-per-week'].quantile(0.75)+1.5*IQR
print('Lower_limit:',Lower)
print('Upper_limit:',Upper)
```

```
Lower_limit: 32.5
Upper_limit: 52.5
```

```
[54]: #getting rid of outliers
df['hours-per-week']=np.where(df['hours-per-week']>Upper,Upper,np.
    ↪where(df['hours-per-week']<Lower,Lower,df['hours-per-week']))
```

```
[55]: #boxplot must not show outliers now
sns.boxplot(x='hours-per-week', data=df);
```



## 10 Separating target variable(Dependent) from Indeendent vari- ables

```
[56]: x=df.iloc[:, :-1]
      y=df['income']
```

```
[57]: # We have Imbalanced Data and we have to do sampling to avoid this problem
      # we have two method so for 1] Under Sampling, 2] Oversampling
      # We will go for Oversampling.
      !pip install -U imbalanced-learn
      from imblearn.over_sampling import RandomOverSampler
      ros=RandomOverSampler()
      x_sample,y_sample=ros.fit_resample(x,y)
```

```
Requirement already satisfied: imbalanced-learn in
c:\users\ps4z\anaconda3\lib\site-packages (0.10.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in
c:\users\ps4z\anaconda3\lib\site-packages (from imbalanced-learn) (2.2.0)
Requirement already satisfied: scipy>=1.3.2 in c:\users\ps4z\anaconda3\lib\site-
packages (from imbalanced-learn) (1.9.1)
```

```
Requirement already satisfied: joblib>=1.1.1 in
c:\users\ps4z\anaconda3\lib\site-packages (from imbalanced-learn) (1.2.0)
Requirement already satisfied: numpy>=1.17.3 in
c:\users\ps4z\anaconda3\lib\site-packages (from imbalanced-learn) (1.21.5)
Requirement already satisfied: scikit-learn>=1.0.2 in
c:\users\ps4z\anaconda3\lib\site-packages (from imbalanced-learn) (1.0.2)
```

```
[58]: #checking our independent variable data
x.head()
```

```
[58]:      age      workclass  fnlwgt  education  education-num  \
0  39.0      State-gov   77516.0    Bachelors             13
1  50.0  Self-emp-not-inc   83311.0    Bachelors             13
2  38.0      Private   215646.0      HS-grad              9
3  53.0      Private   234721.0        11th              7
4  28.0      Private   338409.0    Bachelors             13

      marital-status      occupation  relationship   race   sex  \
0      Never-married      Adm-clerical  Not-in-family  White  Male
1  Married-civ-spouse  Exec-managerial      Husband  White  Male
2      Divorced  Handlers-cleaners  Not-in-family  White  Male
3  Married-civ-spouse  Handlers-cleaners      Husband  Black  Male
4  Married-civ-spouse  Prof-specialty      Wife  Black  Female

      capital-gain  capital-loss  hours-per-week  native-country
0          2174           0         40.0  United-States
1           0           0         32.5  United-States
2           0           0         40.0  United-States
3           0           0         40.0  United-States
4           0           0         40.0         Cuba
```

```
[59]: #checking oit dependent variable data
y.head()
```

```
[59]: 0    0
      1    0
      2    0
      3    0
      4    0
      Name: income, dtype: int64
```

## 11 Label encoding on the categorical features

```
[60]: #Feature Engineering
from sklearn.preprocessing import LabelEncoder
labelEncoder=LabelEncoder()
```

```
[61]: #fit and transform
x[cat_feat]=x[cat_feat].apply(LabelEncoder().fit_transform)
```

```
[62]: #checking independent variable data
x.head()
```

```
[62]:
```

	age	workclass	fnlwgt	education	education-num	marital-status	\
0	39.0	6	77516.0	9	13		4
1	50.0	5	83311.0	9	13		2
2	38.0	3	215646.0	11	9		0
3	53.0	3	234721.0	1	7		2
4	28.0	3	338409.0	9	13		2

	occupation	relationship	race	sex	capital-gain	capital-loss	\
0	0	1	4	1	2174	0	
1	3	0	4	1	0	0	
2	5	1	4	1	0	0	
3	5	0	2	1	0	0	
4	9	5	2	0	0	0	

	hours-per-week	native-country
0	40.0	38
1	32.5	38
2	40.0	38
3	40.0	38
4	40.0	4

```
[63]: #checking target variable data
y.head()
```

```
[63]: 0    0
      1    0
      2    0
      3    0
      4    0
      Name: income, dtype: int64
```

```
[64]: #Hyperparameter Tuning
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,
↪random_state=10)
```

```
[65]: #getting shape of training data of independent variables
x_train.shape
```

```
[65]: (21799, 14)
```

```
[66]: #getting shape of training data of target variable
y_train.shape
```

```
[66]: (21799,)
```

```
[67]: #getting shape of testing data of independent variables
x_test.shape
```

```
[67]: (10738, 14)
```

```
[68]: #getting shape of testing data of target variable
y_test.shape
```

```
[68]: (10738,)
```

## 12 Decision Tree Classifier

```
[69]: #Decision Tree Model
from sklearn.tree import DecisionTreeClassifier
DT=DecisionTreeClassifier()
```

```
[70]: #fitting the model
DT.fit(x_train,y_train)
```

```
[70]: DecisionTreeClassifier()
```

```
[71]: #learning the scores
DT.score(x_test,y_test)
```

```
[71]: 0.8097411063512758
```

```
[72]: #predecting target variable
DTpred=DT.predict(x_test)
```

```
[73]: #Model Evaluation
from sklearn.metrics import accuracy_score
DTacc=accuracy_score(y_test,DTpred)
report=[]
report.append(['Decision Tree', DTacc])
DTacc
```

```
[73]: 0.8097411063512758
```



## 13 Hyperparameter Tuning

```
[74]: #Hyper-Tuning
from sklearn.model_selection import GridSearchCV
grid_parameter={
    'criterion':['gini', 'entropy'],
    'max_depth':range(2,32,1),
    'min_samples_leaf':range(1,10,1),
    'min_samples_split':range(2,10,1),
    'splitter':['best','random']
}
```

```
[75]: DTgrid=GridSearchCV(estimator=DT, param_grid=grid_parameter, cv=3, n_jobs=-1)
```

```
[76]: DTgrid.fit(x_train, y_train)
```

```
[76]: GridSearchCV(cv=3, estimator=DecisionTreeClassifier(), n_jobs=-1,
                  param_grid={'criterion': ['gini', 'entropy'],
                              'max_depth': range(2, 32),
                              'min_samples_leaf': range(1, 10),
                              'min_samples_split': range(2, 10),
                              'splitter': ['best', 'random']})
```

```
[77]: DTgrid.best_params_
```

```
[77]: {'criterion': 'gini',
      'max_depth': 6,
      'min_samples_leaf': 9,
      'min_samples_split': 2,
      'splitter': 'best'}
```

```
[78]: DTbest_para=DecisionTreeClassifier(criterion = "gini", max_depth= 8 ,
    ↪ min_samples_leaf= 9, min_samples_split= 2 ,
    ↪ splitter= "best")
```

```
[79]: DTbest_para.fit(x_train,y_train)
```

```
[79]: DecisionTreeClassifier(max_depth=8, min_samples_leaf=9)
```

```
[80]: DTbest_para_pred2 = DTbest_para.predict(x_test)
```

```
[81]: print("Accuracy Before Hyper-parameter tuning:",accuracy_score(y_test,DTpred))
      print("Accuracy after Hyper-parameter tuning:
      ↪",accuracy_score(y_test,DTbest_para_pred2))
```

Accuracy Before Hyper-parameter tuning: 0.8097411063512758  
Accuracy after Hyper-parameter tuning: 0.8573291115663997

```
[82]: import pickle
      filename='income_pred_model'
      pickle.dump(DT, open(filename, 'wb'))
```

```
[83]: loaded_model=pickle.load(open(filename, 'rb'))
```

```
[84]: loaded_model.predict(x_test)
```

```
[84]: array([0, 0, 1, ..., 1, 0, 1], dtype=int64)
```

```
[85]: df
```

```
[85]:
```

	age	workclass	fnlwgt	education	education-num	\
0	39.0	State-gov	77516.0	Bachelors	13	
1	50.0	Self-emp-not-inc	83311.0	Bachelors	13	
2	38.0	Private	215646.0	HS-grad	9	
3	53.0	Private	234721.0	11th	7	
4	28.0	Private	338409.0	Bachelors	13	
...	...	...	...	...	...	
32556	27.0	Private	257302.0	Assoc-acdm	12	
32557	40.0	Private	154374.0	HS-grad	9	
32558	58.0	Private	151910.0	HS-grad	9	
32559	22.0	Private	201490.0	HS-grad	9	
32560	52.0	Self-emp-inc	287927.0	HS-grad	9	

	marital-status	occupation	relationship	race	\
0	Never-married	Adm-clerical	Not-in-family	White	
1	Married-civ-spouse	Exec-managerial	Husband	White	
2	Divorced	Handlers-cleaners	Not-in-family	White	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	
4	Married-civ-spouse	Prof-specialty	Wife	Black	
...	...	...	...	...	
32556	Married-civ-spouse	Tech-support	Wife	White	
32557	Married-civ-spouse	Machine-op-inspct	Husband	White	
32558	Widowed	Adm-clerical	Unmarried	White	
32559	Never-married	Adm-clerical	Own-child	White	
32560	Married-civ-spouse	Exec-managerial	Wife	White	

	sex	capital-gain	capital-loss	hours-per-week	native-country	\
0	Male	2174	0	40.0	United-States	
1	Male	0	0	32.5	United-States	
2	Male	0	0	40.0	United-States	
3	Male	0	0	40.0	United-States	
4	Female	0	0	40.0	Cuba	
...	...	...	...	...	...	
32556	Female	0	0	38.0	United-States	
32557	Male	0	0	40.0	United-States	

32558	Female	0	0	40.0	United-States
32559	Male	0	0	32.5	United-States
32560	Female	15024	0	40.0	United-States

	income
0	0
1	0
2	0
3	0
4	0
...	...
32556	0
32557	1
32558	0
32559	0
32560	1

[32537 rows x 15 columns]

[86]: x

[86]:

	age	workclass	fnlwgt	education	education-num	marital-status	\
0	39.0	6	77516.0	9	13		4
1	50.0	5	83311.0	9	13		2
2	38.0	3	215646.0	11	9		0
3	53.0	3	234721.0	1	7		2
4	28.0	3	338409.0	9	13		2
...	...	...	...	...	...	...	...
32556	27.0	3	257302.0	7	12		2
32557	40.0	3	154374.0	11	9		2
32558	58.0	3	151910.0	11	9		6
32559	22.0	3	201490.0	11	9		4
32560	52.0	4	287927.0	11	9		2

	occupation	relationship	race	sex	capital-gain	capital-loss	\
0	0	1	4	1	2174	0	
1	3	0	4	1	0	0	
2	5	1	4	1	0	0	
3	5	0	2	1	0	0	
4	9	5	2	0	0	0	
...	...	...	...	...	...	...	...
32556	12	5	4	0	0	0	
32557	6	0	4	1	0	0	
32558	0	4	4	0	0	0	
32559	0	3	4	1	0	0	
32560	3	5	4	0	15024	0	

	hours-per-week	native-country
0	40.0	38
1	32.5	38
2	40.0	38
3	40.0	38
4	40.0	4
...	...	...
32556	38.0	38
32557	40.0	38
32558	40.0	38
32559	32.5	38
32560	40.0	38

[32537 rows x 14 columns]